

A COMPANION GUIDE

The LLM Jargon Glossary

Every term from “Large Language Models: The Jargon You Need to Know” — explained in plain, simple words.

Each term gets a one-line, plain-English summary, followed by a short, detailed explanation with an everyday example — good for revision, interviews, or the next time this vocabulary shows up in a job description.

Sheryians AI School

How to Use This Guide

This is the companion guide to the video “Large Language Models: The Jargon You Need to Know.” Every term covered in the video is explained here twice: first in one plain, simple sentence, then again in a short, detailed paragraph — in everyday words, with an example wherever it helps. Keep this handy for interviews, revision, or whenever a new AI term shows up in a job description.

WHAT'S INSIDE

How We Got Here

Transformer

What Is an LLM

Large Language Model (LLM)

The Model

Parameters, Model Size, Foundation Model

Text & Context

Tokens, Context Window, Prompt

Controlling Output

Temperature, System Prompt, Inference

Representing Information

Embeddings

Training

Training, Fine-Tuning, RLHF

Modern LLM Concepts

Hallucination, RAG (Retrieval-Augmented Generation), Quantization, LoRA / PEFT, Multimodal Models

Applications

Interview Cheat Sheet

Quick phrase-to-term lookup

How We Got Here

Before today's chatbots existed, computers learned in a much simpler way: look at examples, find patterns, make a guess. Over the years, those models kept getting bigger and more capable — until one invention changed everything.

THE BIG INVENTION

Transformer

SIMPLY PUT

A new computer design that lets a model pay attention to the most important words in a sentence.

Before 2017, computers mostly read text word by word, in order, and struggled to remember what happened earlier in a long sentence. The Transformer changed that. It introduced an idea called “attention,” which lets a model look at an entire sentence at once and figure out which words matter most to understand the meaning. For example, in “The dog chased the ball because it was fast,” a Transformer can work out that “it” means the dog, not the ball. This single idea is the invention that made today's LLMs possible.

What Is an LLM

THE BASICS

Large Language Model (LLM)

SIMPLY PUT

A computer program that read a huge amount of text and learned how words fit together.

An LLM is trained on an enormous pile of text — books, websites, articles, code, conversations. It doesn't “know” facts the way a person does. Instead, it's extremely good at predicting what word (technically, what token) should come next, based on everything it has read before. Do that well enough, at a big enough scale, and the result can hold conversations, answer questions, write stories, and even write computer code — all from that one simple skill of “guess the next word.”

The Model

INSIDE THE BRAIN

Parameters

SIMPLY PUT

The tiny, learned settings inside a model's brain that help it make predictions.

Think of parameters as millions (or billions) of tiny dials inside the model, each one slightly adjusted during training so the model gets better at predicting text. When you hear a model has “70 billion parameters,” that means it has 70 billion of these learned dials. Generally, more parameters lets a model learn more complex patterns — but it also needs more computing power to run.

HOW BIG IS IT

Model Size

SIMPLY PUT

How many parameters a model has — usually described as small, medium, or large.

Models are often labeled by size: “small” might mean around 1 billion parameters, “medium” around 7 billion, and “large” 70 billion or more. A bigger model is usually more capable, similar to how a bigger brain can hold more information. But bigger also means slower and more expensive to run — so companies often pick the smallest model that is still good enough for the job at hand.

THE BASE LAYER

Foundation Model

SIMPLY PUT

One big, general-purpose model that other, more specialized tools get built on top of.

A foundation model is trained on a huge, broad range of data so it's generally capable at many things. It's called a “foundation” because it acts like a base layer — companies take it and build more specific tools on top, like a customer-support chatbot or a coding assistant. Think of it like a big block of clay: useful on its own, but also ready to be shaped into something more specific.

Text & Context

CHOPPING UP TEXT

Tokens

SIMPLY PUT

The small word-pieces a model breaks text into before it can process it.

A token isn't always a whole word. Sometimes it's part of one — for example, “unbelievable” might get split into “un” and “believable.” The model reads and writes one token at a time, predicting the next token based on everything that came before it. This is also the unit that companies use for pricing — when you hear about “cost per token” or a “token limit,” this is what's being counted.

SHORT-TERM MEMORY

Context Window

SIMPLY PUT

How much text a model can remember and consider at one time.

The context window is the model's short-term memory. If a conversation or document goes on longer than the context window, the model starts to “forget” the earliest parts. Context windows are measured in tokens — common sizes are 8,000 (8K), 128,000 (128K), or even 1 million (1M) tokens. A bigger context window means the model can keep track of longer conversations, documents, or codebases at once.

JUST ASK

Prompt

SIMPLY PUT

The message or instruction you type to the model — your question or request.

Everything a model produces is a response to a prompt. A vague prompt tends to get a vague answer; a clear, specific prompt tends to get a clear, useful one. This is exactly where the term “prompt engineering” comes from — the skill of writing prompts that reliably get the result you actually want.

Controlling Output

SAFE VS. CREATIVE

Temperature

SIMPLY PUT

A setting that controls how safe or how random a model's answers are.

A low temperature makes the model stick to its single most likely, safest answer every time — good for factual tasks where you want consistency. A high temperature lets the model take more creative risks, producing more varied and surprising answers — good for brainstorming or creative writing. It's usually set as a number between 0 and 1 (sometimes higher).

THE HIDDEN NOTE

System Prompt

SIMPLY PUT

A hidden instruction that tells the model how to behave, set before the conversation begins.

A system prompt is like a note the developers leave for the model in advance — things like “You are a friendly, concise tutor” or “Only answer questions about cooking.” The user usually never sees this note, but it quietly shapes the model's tone and behavior throughout the entire conversation.

PUTTING IT TO WORK

Inference

SIMPLY PUT

Using an already-trained model to actually generate an answer.

Training happens once, and it's expensive. Inference happens every single time someone sends the model a message — it's the model putting everything it learned to use, right now, to answer you. When people talk about “inference cost” or “inference speed,” they mean how much it costs, or how fast it is, to get an answer out of the model in everyday, real-world use.

Representing Information

WORDS AS NUMBERS

Embeddings

SIMPLY PUT

A way of turning words into numbers so a computer can measure how similar they are.

An embedding turns a word, sentence, or even an image into a long list of numbers — a point in space. Words with similar meanings end up as points that are close together, a bit like plotting them on a map. “King” and “queen” would land near each other, while “car” would land far away. This is how a model understands that words are related in meaning, not just in spelling.

Training

PRACTICE, PRACTICE, PRACTICE

Training

SIMPLY PUT

The process of teaching a model by having it practice on huge amounts of example data.

During training, the model makes a guess, checks how wrong it was, and slightly adjusts its internal parameters to do better next time. This repeats over and over — often billions of times — until the model gets good at predicting text. It's a bit like practicing free throws: shoot, see if you missed, adjust your aim, and repeat.

EXTRA PRACTICE

Fine-Tuning

SIMPLY PUT

Giving an already-trained model extra, focused practice on one specific topic.

Instead of training a brand-new model from scratch, fine-tuning takes a general-purpose foundation model and teaches it to be really good at one particular thing — like answering legal questions, or writing in a specific company's tone. It's a much faster and cheaper way to specialize a model than starting over from zero.

LEARNING FROM PEOPLE

RLHF**SIMPLY PUT**

A training method where real people rate a model's answers to teach it to be more helpful.

RLHF stands for Reinforcement Learning from Human Feedback. The model produces a few different responses to the same question, and a human rates or ranks them — this one is better, that one is worse. The model then learns to produce more of what people rated highly. This is a big reason modern AI assistants feel more helpful and polite than early, purely text-predicting models.

Modern LLM Concepts

CONFIDENTLY WRONG

Hallucination**SIMPLY PUT**

When a model gives an answer that sounds right but is actually made up.

A hallucination happens because the model is built to predict likely-sounding text, not to look up guaranteed facts. It might invent a fake statistic, a fake quote, or even a book that doesn't exist — all while sounding completely confident. This is exactly why it's important to double-check any important fact a model gives you.

LOOKING IT UP FIRST

RAG (Retrieval-Augmented Generation)**SIMPLY PUT**

Letting a model look up real information before it answers, instead of only guessing from memory.

With RAG, before the model writes its answer, it first retrieves relevant information from a trusted source — like a company's documents or a search engine. It then uses that real information to write its response. This helps reduce hallucinations, because the model is “grounded” in real, current data instead of relying purely on what it memorized during training.

SHRINKING THE BRAIN

Quantization

SIMPLY PUT

Making a model smaller and faster by rounding its internal numbers to take up less space.

A model's internal numbers can be stored with more or less precision. Full-precision numbers (like FP32) take up a lot of memory; quantized versions (like INT8 or INT4) use shorter, rounded versions of those same numbers. The model loses a tiny bit of accuracy, but becomes dramatically smaller, faster, and cheaper to run — which matters a lot for running models on phones or smaller devices.

THE EFFICIENT SHORTCUT

LoRA / PEFT

SIMPLY PUT

A shortcut that fine-tunes a model by training a small add-on part instead of the whole brain.

LoRA (Low-Rank Adaptation) is part of a family of techniques called PEFT (Parameter-Efficient Fine-Tuning). Instead of retraining all of a model's billions of parameters, LoRA adds a small, separate set of new, trainable parameters — and only trains those. This makes customizing a model dramatically cheaper and faster, since you're training a tiny add-on instead of the entire brain.

MORE THAN WORDS

Multimodal Models

SIMPLY PUT

A model that understands more than text — it can also work with images, audio, or video.

“Modal” refers to a type, or mode, of information — so “multimodal” just means “many types.” A multimodal model might let you show it a photo and ask “what's in this picture?” or give it an audio clip and ask it to summarize what was said. The same underlying model handles all of these input types, not just plain text.

Applications

Here's the part that ties it all together: the same underlying LLM can power a huge range of different tools — chatbots, coding assistants, translators, summarizers, search tools, and content generators. The model itself usually doesn't change between these; what changes is how it's prompted, and what extra tools (like RAG) are wrapped around it.

Interview Cheat Sheet

Quick clues — hear one of these phrases, know which word it means.

IF YOU HEAR THIS...	THEY MEAN...
"70B model"	Parameters
"128K context"	Context Window
"The model is hallucinating"	Hallucination
"Ground the model"	RAG
"Make the model smaller"	Quantization
"Customize the model"	Fine-Tuning / LoRA
"Use the trained model"	Inference
"The model accepts image + text"	Multimodal
"It's read the whole internet"	Large Language Model (LLM)
"It can only remember so much"	Context Window
"Turn the creativity up"	Temperature
"It's cheap to run on a phone"	Quantization

That's the full vocabulary from the video. Keep this glossary open next to your notes, or skim it once before an interview — every term here is one you'll keep running into in AI job descriptions, docs, and interviews.